

ASPP CUDA programming coursework

Please see Learn for submission deadlines.

Remember what you submit must be your own work. Anything not your own work which you have accessed should be correctly referenced and cited. You must not share this assessment's source code nor your solutions. Use of generative AI tools, in line with School policy, is not permitted.

Please see further information and guidance from the School of Informatics Academic Misconduct Officer: <https://web.inf.ed.ac.uk/infweb/admin/policies/academic-misconduct>

Summary

Your goal here is to take a naïve, serial code that runs a percolation model, and make this run correctly and efficiently using a single GPU on Cirrus (i.e. one NVIDIA V100).

You must also prepare a brief report (maximum 1 page) explaining, with reference to the algorithm and hardware, how the changes you have made to the code achieve this performance.

As part of marking, your code will be compiled and run on Cirrus using the `gpu` partition, so please use that for any tuning and profiling you do.

It is **ABSOLUTELY ESSENTIAL** that you name your submitted source file as `$EXAMNUMBER.cu`, for example `B654321.cu` and this must match the submitted report. Your code will be compiled and run automatically. Deviating from this is likely to result in a mark of zero for correctness and performance.

Expected time to complete: 5 hours

Problem description

So called “percolation models” are used as simplified models of various systems, including forest fires, disease propagation, and flow in porous media. Here we focus on the latter, asking: given a random lattice 2D material with a porosity p (i.e. the material is approximated as a grid of equal sized squares, each being empty with independent probability p), do the pore spaces connect?

This can be done by first labelling each non-solid cell with a unique number, then iteratively updating this value to the maximum of the labels at that cell and its four immediate neighbours, i.e.

```
new_label[i,j] = max(  
    label[ i,j+1],  
    label[i-1, j], label[ i, j], label[i+1, j]  
    label[ i,j-1]  
)
```

To know when to finish, the algorithm needs to know when this process has converged, i.e. when the total number of changes across the whole grid is zero.

The supplied code applies this, in serial, on the CPU. The driver code runs this once on the CPU to compute a reference solution, then a number of times using the GPU implementation (in the initial code, this is simply a copy of the CPU implementation with an altered function name), before validating the solution, printing timing statistics, and writing the a image of the output.




Figure 1: Sample output image

It accepts the following flags:

- M `<integer>` - horizontal size of grid (default 512)
- N `<integer>` - vertical size of grid (default 512)
- S `<integer>` - random seed (default 1234)

- r <integer> - number of repeats for benchmarking (default 3)
- p <float> - target porosity between 0 and 1 (default 0.4)
- o <path> - file name to write output PNG image (default `test.png`)

Set up

On Cirrus:

1. Clone the code repository (recalling that only the work filesystem is accessible from the compute nodes): `git clone /home/z04/shared/aspp/cw1`
2. Load the required modules: `module load cmake gcc/10.2.0 nvidia/nvhpc/24.5`
3. Configure the code. For the marking runs, I will use a build type of **Release**, but you may prefer to build with some debug info while testing. You can put the build directory where you like, but the included scripts expect it to be in `$repo/build`

```
cmake -S src -B build \
-DMAKE_BUILD_TYPE=RelWithDebInfo \
-DMAKE_CUDA_ARCHITECTURES=70-real
```

4. Compile: `cmake --build build -j 4`. This will produce an executable `test`.
5. Running the *unmodified* code is possible on the login nodes. If you run with no options (just `build/test`) the produced image should be identical to the `default_output.png` in the root of the repository.
6. I have provided several template batch scripts, one (`run.sh`) simply runs the executable while the other two profile it.

Requirements for your code

You must adapt the code so that it runs on one GPU, correctly and with good performance.

Your code will form part of the submission and is worth 50% of the marks for this assignment. You may only submit your changed version of `perc_gpu.cu`. It is **ESSENTIAL** that you name the submitted file `YOUR_EXAM_NUMBER.cu` (the one that starts with a “B”). Calling it anything else may make the test fail and make it impossible to match up with your report. This submitted file will be added to a copy of the supplied repository for benchmarking.

Your code must compile (with only the modules given above loaded) by running `cmake` as above (with `MAKE_BUILD_TYPE=Release`) and run with the existing command line options. Note that a code that does not compile is, by definition

incorrect, and that a code that produces incorrect results scores no points for performance.

You may use functions and types from the standard library and base CUDA library, but from no other sources.

To be sure your code works, please test it on Cirrus and submit that version, correctly named!

Clarity: your modifications will be marked for usual good practice in programming. Are variables/functions sensibly named? Code well formatted? Are comments present, where necessary, to explain what is not obvious?

Correctness: the benchmark code compares the GPU results to those from the (unmodified) serial CPU version. I will compile and run this against a range of problem sizes and porosities.

Performance: I will use several problems of different sizes and porosities, taking the best performance run in each case. Sizes will be between 128 and 16,384 inclusive along *each* dimension.

Requirements for report

This must be a PDF with a maximum of one A4 page and text at 10pt or greater.

1. Why is this problem more complex to solve on a GPU than a single core? (2-3 sentences) [10 %]
2. Explain the choices that you have made to get good performance and how you arrived at them, with reference to *both* the hardware and the problem to solve. You may choose to refer to your code (e.g. see lines 90–95). [30 %]
3. What level of performance might be achievable and how close have you got to it? [10%]

Submission

Please see the instructions on Learn for full details, but you will have to submit the code and report to separate queues.

It is VITAL that you use your exam number for both the code and report filenames so we can match them up! I.e., the report should be B123456.pdf and the code B123456.cu.

Hints

You have been provided the code in a version controlled Git repository. I suggest that you make use of it to track your progress and you will be able to tell if you have inadvertently modified files that you should not with `git status`.

I highly recommend starting with a look at the existing code and thinking about point 1 from the report. Consider how you will decompose the index space between blocks and threads and how you will decide when to terminate.

You should test with a variety of problem sizes.

You may wish to get interactive access to a GPU if you are doing debugging/performance tuning. Use this SLURM command:

```
srun --partition=gpu --qos=short --gres=gpu:1 \  
    --time=0:10:0 --account=$BUDGETCODE --pty /usr/bin/bash --login
```

(replacing `$BUDGETCODE` with your budget).

I note that running extra iterations should not change the result, but that this will take more time.